

Distributed Systems (Assignment-3):

Problem:

Design and implement a Remote Procedure Call framework to simulate the following:

- Implement a single server architecture with support for multiple clients
- The server maintains a list of graphs each associated with a distinct identifier.
- Clients can request to add a new graph, update an existing graph and query for the total weight of the minimum weight spanning tree of a given graph.
- Clients can request to add a new graph using '**add_graph** <graph_identifier> n'. This command will add a new graph on the server with the identifier 'graph_identifier' and 'n' number of nodes. The graph identifier is a string with a maximum length of 10 and it won't already exist. 'n' will be in the range: $1 \leq n \leq 100,000$.
- Clients can request to add a new edge in a graph using '**add_edge** <graph_identifier> <u> <v> <w>'. This will add an undirected edge between the nodes **u** and **v** with weight **w**. **u** and **v** are the node numbers of the endpoints of the edge such that $1 \leq u, v \leq n$ and $0 \leq w \leq 10,000$. **n** is the number of nodes in the specified graph. A graph with identifier graph_identifier will already exist. There can be multiple edges and self-loops added to the graph.
- Clients can request for the total weight of the minimum weight spanning tree in a graph from the server using '**get_mst** <graph_identifier>'. The client will print the solution the server returns. In case the graph does not have a spanning tree, -1 should be printed. A graph with identifier graph_identifier will already exist.
- All values should fit in 32-bit signed integers.
- The server should be able to handle multiple clients simultaneously and should also work with clients on other machines.
- You are free to use any algorithm for MST

Solution Approach

- Programming language for Client and Server - C++
- Using Apache **'thrift'** package for Remote Procedure Calls.
- Create an **interface.thrift** IDL file for thrift compiler to generate the RPC service code. This file will contain the interfaces that will be exposed for the client to perform a remote procedure call to the server. Given below is the IDL file used for current assignment

```
service GraphService
{
    oneway void addGraph(1:string graphId,
                        2:i32 nodeCount)

    i32 addEdge(1:string graphId,
                2:i32 srcNode,
                3:i32 dstNode,
                4:i32 edgeWeight)

    i32 getMSTWeight(1:string graphId)
}
```

On Client Side:

Implement a menu-driven approach for performing all the above operations. A typical menu implemented in the client is as shown below:

```
| Add a Graph      ( Command: add_graph <graph_id_string> N ) |
| Add an Edge     ( Command: add_edge <graph_id_string> <u> <v> <w> ) |
| Get MST Weight  ( Command: get_mst <graph_id_string> ) |
|_____|
|                                     "-1" to EXIT |
|_____|
Command: _
```

On Server Side:

- Use a `std::map<key, value>` as a graph container. The 'key' would be a string to identify the Graph and 'value' would be a pointer to the Graph.
- Implement a 'Graph' class encapsulating all the interfaces that can be called from the server. Use Kruskal's algorithm for computing the Minimum Spanning Tree for an undirected graph, using *Union-Find* data structure for better performance.

Environment Details

Operating System: *Mac OSX BigSur 11.2.2*
Package Installer : *Homebrew*
Compiler : *clang++*
RPC Package : *Apache Thrift*

Installation Instructions:

- Install CommandLineTools and Homebrew package manager on Mac OSX. Refer [this](#) link for detailed installation steps.
- The C++ compiler 'clang++' would come along with the CommandLineTools package.
- Using Homebrew package installer install Apache 'thrift' as follows:

```
brew install thrift
```

- Make a note of the include and the dynamic library paths in the Apache 'thrift' installation to update the Makefile. A typical include and library paths would look like:
/usr/local/Cellar/thrift/0.14.0/include
/usr/local/Cellar/thrift/0.14.0/lib
- Update the INCLUDE and LD_LIBRARY_PATH variables in the Makefile accordingly, using the above paths from your installation.

Source Tree

Given below is the source tree for the current assignment.

```
|— DS_Assignment3.pdf
|— Makefile
|— README.md
|— bin/
|— obj/
└─ src/
    |— interface.thrift
    |— client/
    |   |— Client.cpp
    |   |— ClientUtils.cpp
    |   └─ ClientUtils.h
    |— server/
    |   |— Graph.cpp
    |   |— Graph.h
    |   |— GraphServiceHandler.cpp
    |   |— GraphServiceHandler.h
    |   |— GraphUtils.cpp
    |   |— GraphUtils.h
    |   └─ Server.cpp
    └─ service/
        |— GraphService.cpp
        |— GraphService.h
        |— GraphService_server.skeleton.cpp
        └─ interface_types.h
```

- 'src/service/' : Directory containing auto generated files from Apache 'thrift' compiler.
- 'src/client/' : Directory containing all the client code & utility files.
- 'src/server' : Directory containing all the server code & utility files.
- Client.cpp : Code to instantiating a client.
- ClientUtils.* : Utility functions for displaying the menu and processing the input commands.
- Server.cpp : Code to instantiating a server.
- Graph.* : A class implementing all the interfaces to create the graph and algorithm for computing Minimum Spanning Tree weight.
- GraphUtils.* : Implementation of Union-Find data structure, used while computing MST.
- GraphServiceHandler.* : A service handler class derived from the RPC service files created by Apache 'thrift' compiler. Used by the server code to execute the RPC function calls.
- Run the command '**make clean; make**' from the parent directory to create the *client* and *server* executable binaries.

Building the *client* and *server* binaries

After updating the Makefile, build the client and server binaries as follows:

```
20173071$ make clean; make
Cleaning the object files and binaries.
rm -f core ./bin/* ./obj/*

Generating interface files.
mkdir -p ./src/service
thrift -out ./src/service --gen cpp ./src/interface.thrift

clang++ -Wall -std=c++11 -I./ -I/usr/local/Cellar/thrift/0.14.0/include -c src/client/Client.cpp -o obj/Client.o
Compiled src/client/Client.cpp successfully.
clang++ -Wall -std=c++11 -I./ -I/usr/local/Cellar/thrift/0.14.0/include -c src/client/ClientUtils.cpp -o obj/ClientUtils.o
Compiled src/client/ClientUtils.cpp successfully.
clang++ -Wall -std=c++11 -I./ -I/usr/local/Cellar/thrift/0.14.0/include -c src/service/GraphService.cpp -o obj/GraphService.o
Compiled src/service/GraphService.cpp successfully.
clang++ -lthrift ./obj/Client.o ./obj/ClientUtils.o ./obj/GraphService.o -o ./bin/client
Linking client objects complete.
To start the "Client" run --> ./bin/client
```

```
clang++ -Wall -std=c++11 -I./ -I/usr/local/Cellar/thrift/0.14.0/include -c src/server/Server.cpp -o obj/Server.o
Compiled src/server/Server.cpp successfully.
clang++ -Wall -std=c++11 -I./ -I/usr/local/Cellar/thrift/0.14.0/include -c src/server/GraphServiceHandler.cpp -o obj/GraphServiceHandler.o
Compiled src/server/GraphServiceHandler.cpp successfully.
clang++ -Wall -std=c++11 -I./ -I/usr/local/Cellar/thrift/0.14.0/include -c src/server/Graph.cpp -o obj/Graph.o
Compiled src/server/Graph.cpp successfully.
clang++ -Wall -std=c++11 -I./ -I/usr/local/Cellar/thrift/0.14.0/include -c src/server/GraphUtils.cpp -o obj/GraphUtils.o
Compiled src/server/GraphUtils.cpp successfully.
clang++ -lthrift ./obj/Server.o ./obj/GraphServiceHandler.o ./obj/Graph.o ./obj/GraphUtils.o ./obj/GraphService.o -o ./bin/server
Linking server objects complete.
To start the "Server" run --> ./bin/server

20173071$
```

Execution Run

- Instantiate ***./bin/server*** from one terminal
- Instantiate ***./bin/client*** from another terminal. Multiple client instantiations would require multiple terminals.
- Follow the menu in the client terminals to create the undirected graph and compute the MST weight.
- The log traces are categorized as below:

```
[INFO] -> Informational traces
[RPCQ] -> RPC Request traces
[RPCR] -> RPC Response traces
[ERROR] -> Error traces
```

Additional Features

- After the graph is created with the initial number of nodes provided, we can add additional edges with new nodes (not part of the initial graph) and request for the MST. Dynamically the server would re-configure the graph with additional nodes and edges and compute the MST.
- Useful when you try to test where MST wouldn't exist. I added a connected graph with 4 nodes and computed the MST. Then I added another disconnected edge (which would add additional 2 nodes and 1 edge into the initial graph of 4 nodes). MST weight should return -1.
- Now add another edge between one of the node from original graph and any one of the node in the newly added disconnected edge. Now the new graph is connected with 6 nodes. MST should exist and hence should return a value not equal to -1.