

# Lecture 4

15<sup>th</sup> January 2020

Team 6: Amrit Sahai(2019201009), Hetav Panchani(2018101016), Devansh Gupta(201711100)

## ***PRESENTATION SCRIBE***

***TOPIC: How vector clock deals with out of order messages ?***

### ***CHALLENGES:***

1. Receive the messages in the order they were sent.
2. If a message with higher timestamp is received before a lower timestamp message, the protocol should detect that there is a message yet to arrive which was sent earlier.
3. When such a case is present, the current received message may be stored in a buffer until the earlier sent message is received first.

### **PROTOCOLS FOR HANDLING OUT OF ORDER MESSAGES**

1. Birman-Schiper-Stephenson Protocol
2. Schiper-Eggli-Sandoz Protocol

### **Birman-Schiper-Stephenson Protocol**

The goal of this protocol is to preserve ordering in the sending of messages. For example, if  $send(m_1) \rightarrow send(m_2)$ , then for all processes that receive both  $m_1$  and  $m_2$ ,  $receive(m_1) \rightarrow receive(m_2)$ . The basic idea is that  $m_2$  is not given to the process until  $m_1$  is given. This means a buffer is needed for pending deliveries. Also, each message has an associated vector that contains information for the recipient to determine if another message preceded it. Also, we shall assume all messages are broadcast. Clocks are updated only when messages are sent.

### **Notation**

- $N$  processes
- $P_i$  process

- $C_i$  vector clock associated with process  $P_i$ ;  $j$ th element is  $C_i[j]$  and contains  $P_i$ 's latest value for the current time in process  $P_j$
- $t^m$  vector timestamp for message  $m$  (stamped after local clock is incremented)

## Protocol


**$P_i$  sends a message to  $P_j$**

$P_i$  increments  $C_i[i]$  and sets the timestamp  $t^m = C_i[i]$  for message  $m$ .

**$P_j$  receives a message from  $P_i$**

- When  $P_j, j \neq i$ , receives  $m$  with timestamp  $t^m$ , it delays the message's delivery until both
  1.  $C_j[i] = t^m[i] - 1$ ; and
  2. for all  $k \leq n$  and  $k \neq i, C_j[k] \leq t^m[k]$
- When the message is delivered to  $P_j$ , update  $P_j$ 's vector clock
- Check buffered messages to see if any can be delivered.

**EXAMPLE:**



*Here is the protocol applied to the above situation:*

$e31$ :  $P_3$  sends message  $a$ ;  $C_3 = (0, 0, 1)$ ;  $t^a = (0, 0, 1)$

$e21$ :  $P_2$  receives message  $a$ . As  $C_2 = (0, 0, 0)$ ,  $C_2[3] = t^a[3] - 1 = 1 - 1 = 0$  and  $C_2[1] \Rightarrow t^a[1]$  and  $C_2[2] \Rightarrow t^a[2] = 0$ . So the message is accepted, and  $C_2$  is set to  $(0, 0, 1)$

$e11$ :  $P_1$  receives message  $a$ . As  $C_1 = (0, 0, 0)$ ,  $C_1[3] = t^a[3] - 1 = 1 - 1 = 0$  and  $C_1[1] \Rightarrow t^a[1]$  and  $C_1[2] \Rightarrow t^a[2] = 0$ . So the message is accepted, and  $C_1$  is set to  $(0, 0, 1)$

$e22$ :  $P_2$  sends message  $b$ ;  $C_2 = (0, 1, 1)$ ;  $t^b = (0, 1, 1)$

$e12$ :  $P_1$  receives message  $b$ . As  $C_1 = (0, 0, 1)$ ,  $C_1[2] = t^b[2] - 1 = 1 - 1 = 0$  and  $C_1[1] \Rightarrow t^b[1]$  and  $C_1[3] \Rightarrow t^b[3] = 0$ . So the message is accepted, and  $C_1$  is set to  $(0, 1, 1)$

$e32$ :  $P_3$  receives message  $b$ . As  $C_3 = (0, 0, 1)$ ,  $C_3[2] = t^b[2] - 1 = 1 - 1 = 0$  and  $C_3[1] \Rightarrow t^b[1]$  and  $C_3[3] \Rightarrow t^b[3] = 0$ . So the message is accepted, and  $C_3$  is set to  $(0, 1, 1)$

Now, suppose  $t^a$  arrived as event  $e12$ , and  $t^b$  as event  $e11$ . Then the progression of time in  $P_1$  goes like this:


$e11$ :  $P_1$  receives message  $b$ . As  $C_1 = (0, 0, 0)$ ,  $C_1[2] = t^b[2] - 1 = 1 - 1 = 0$  and  $C_1[1] \Rightarrow t^b[1]$ , but  $C_1[3] < t^b[3]$ , so the message is held until another message arrives. The vector clock updating algorithm is not run.

$E12$ :  $P_1$  receives message  $a$ . As  $C_1 = (0, 0, 0)$ ,  $C_1[3] = t^a[3] - 1 = 1 - 1 = 0$ ,  $C_1[1] \Rightarrow t^a[1]$ , and  $C_1[2] \Rightarrow t^a[2]$ . The message is accepted and  $C_1$  is set to  $(0, 0, 1)$ . Now the queue is checked. As  $C_1[2] = t^b[2] - 1 = 1 - 1 = 0$ ,  $C_1[1] \Rightarrow t^b[1]$ , and  $C_1[3] \Rightarrow t^b[3]$ , that message is accepted and  $C_1$  is set to  $(0, 1, 1)$ .

## Schiper-Eggli-Sandoz Protocol

The goal of this protocol is to ensure that messages are given to the receiving processes in order of sending. Unlike the Birman-Schiper-Stephenson protocol, it does not require using broadcast messages. Each message has an associated vector that contains information for the recipient to determine if another message preceded it. Clocks are updated only when messages are sent.

### EXAMPLE:



### Here is the protocol applied to the above situation:

*e31*:  $P_3$  sends message  $a$  to  $P_2$ .  $C_3 = (0, 0, 1)$ ;  $t^a = (0, 0, 1)$ ,  $V^a = (?, ?, ?)$ ;  $V_3 = (?, (0, 0, 1), ?)$

*e21*:  $P_2$  receives message  $a$  from  $P_1$ . As  $V^a[2]$  is uninitialized, the message is accepted.  $V_2$  is set to  $(?, ?, ?)$  and  $C_2$  is set to  $(0, 0, 1)$ .

*e22*:  $P_2$  sends message  $b$  to  $P_1$ .  $C_2 = (0, 1, 1)$ ;  $t^b = (0, 1, 1)$ ,  $V^b = (?, ?, ?)$ ;  $V_2 = ((0, 1, 1), ?, ?)$

*e11*:  $P_1$  sends message  $c$  to  $P_3$ .  $C_1 = (1, 0, 0)$ ;  $t^c = (1, 0, 0)$ ,  $V^c = (?, ?, ?)$ ;  $V_1 = (?, ?, (1, 0, 0))$

*e12*:  $P_1$  receives message  $b$  from  $P_2$ . As  $V^b[1]$  is uninitialized, the message is accepted.  $V_1$  is set to  $(?, ?, ?)$  and  $C_1$  is set to  $(1, 1, 1)$ .

*e32*:  $P_3$  receives message  $c$  from  $P_1$ . As  $V^c[3]$  is uninitialized, the message is accepted.  $V_3$  is set to  $(?, ?, ?)$  and  $C_3$  is set to  $(1, 0, 1)$ .

*e23*:  $P_2$  sends message  $d$  to  $P_1$ .  $C_2 = (0, 2, 1)$ ;  $t^d = (0, 2, 1)$ ,  $V^d = ((0, 1, 1), ?, ?)$ ;  $V_2 = ((0, 2, 1), ?, (0, 0, 1))$

*e13*:  $P_1$  receives message  $d$  from  $P_2$ . As  $V^d[1] < C_1[1]$ , so the message is accepted.  $V_1$  is set to  $((0, 1, 1), ?, ?)$  and  $C_1$  is set to  $(1, 2, 1)$ .

Now, suppose  $t^b$  arrived as event *e13*, and  $t^d$  as event *e12*. Then the progression in  $P_1$  goes like this:

e12:  $P_1$  receives message  $d$  from  $P_2$ . But  $V^d[1] = (0, 1, 1) <_X (1, 0, 0) = C_3$ , so the message is queued for later delivery.

E13:  $P_1$  receives message  $b$  from  $P_2$ . As  $V^b[1]$  is uninitialized, the message is accepted.  $V_1$  is set to  $(?, ?, ?)$  and  $C_1$  is set to  $(1, 1, 1)$ . The message on the queue is now checked. As  $V^d[1] = (0, 1, 1) < (1, 1, 1) = C_1$ , the message is now accepted.  $V_1$  is set to  $((0, 1, 1), ?, ?)$  and  $C_1$  is set to  $(1, 2, 1)$ .

## **LECTURE SCRIBE:**

### **Message ordering:**

- FIFO
- NON-FIFO
- CAUSAL

In the FIFO model, each channel acts as a first-in first-out message queue and, thus, message ordering is preserved by a channel. In the non-FIFO model, a channel acts like a set in which the sender process adds messages and the receiver process removes messages from it in a random order. A system that supports causal delivery of messages satisfies the following property: “for any two messages  $m_{ij}$  and  $m_{kj}$ , if  $\text{send}(m_{ij}) \rightarrow \text{send}(m_{kj})$ , then  $\text{rec}(m_{ij}) \rightarrow \text{rec}(m_{kj})$ ”

Causally ordered delivery of messages implies FIFO message delivery. The causal ordering model is useful in developing distributed algorithms and may simplify the design of algorithms.

### **Matrix Time**

It is the information that a process has about other processes and their information about other processes. In a system of matrix clocks, the time is represented by a set of  $n \times n$  matrices of non-negative integers. A process  $p_i$  maintains a matrix  $mti[1..n, 1..n]$  where,

- $mti[i, i]$  denotes the local logical clock of  $p_i$  and tracks the progress of the computation at process  $p_i$ ;

- $mti[i,j]$  denotes the latest knowledge that process  $p_i$  has about the local logical clock,  $mt[j,j]$ , of process  $p_j$  (note that row,  $mti[i, \cdot]$  is nothing but the vector clock  $v_i[\cdot]$  and exhibits all the properties of vector clocks);
- $mti[j,k]$  represents the knowledge that process  $p_i$  has about the latest knowledge that  $p_j$  has about the local logical clock,  $mt[k,k]$ , of  $p_k$ .

The entire matrix  $mti$  denotes  $p_i$ 's local view of the global logical time. The matrix timestamp of an event is the value of the matrix clock of the process when the event is executed. Process  $p_i$  uses the following rules R1 and R2 to update its clock:

- R1: Before executing an event, process  $p_i$  updates its local logical time as follows:  

$$mti[i,i] := mti[i,i] + d \quad (d > 0)$$
- R2: Each message  $m$  is piggybacked with matrix time  $mt$ . When  $p_i$  receives such a message  $(m, mt)$  from a process  $p_j$ ,  $p_i$  executes the following sequence of actions:
  - update its global logical time as follows:
  - $1 \leq k \leq n$ :  $mti[i,k] := \max(mt[i,k], mt[j,k])$ , (that is, update its row  $mti[i, \cdot]$  with  $p_j$ 's row in the received timestamp,  $mt$ );
    - $1 \leq k, l \leq n$ :  $mti[k,l] := \max(mti[k,l], mt[k,l])$ ;
  - execute R1;
  - deliver message  $m$ .

Q. Is the principal vector always larger than the non-principal vectors? **TRUE**

This is so because the user has the latest information about its processes. When the information from other processes is received, the principle vector is updated just after receiving and hence, it is always the most updated one.

# Global State of a Distributed System

At any instant the state of a process  $P_i$  denoted by  $LS_i$ , is a result of the sequence of all events executed by  $p_i$  upto that instant.

$$SC_{ij} = \{m_{ij} \mid \text{send}(m_{ij}) \in LS_i \text{ and } \text{rec}(m_{ij}) \in LS_j\}$$

To formalize, let  $LS_{x_i}$  denote the local state of process  $P_i$  after the occurrence of all events until event  $e_{x_i}$

- $LS_{0_i}$  is the initial state of  $P_i$

## Conditions for a consistent global state

A global state  $GS = \{U_i, LS_{x_i}, U_j, SC_{y_j}, z_k, j_k\}$  is a consistent global state iff it satisfies the following condition:

$$\forall m_{ij} : \text{send}(m_{ij}) \leq LS_{x_i} \Rightarrow m_{ij} \notin SC_{y_j} \wedge \text{rec}(m_{ij}) \leq LS_{y_j}$$

That is, channel state  $SC_{y_i, z_k}$  and process state  $LS_{z_k}$  must not include any message that process  $p_i$  sent after executing event  $e_{x_i}$

## Homework Questions:

Q1. Taking example of using matrix clock (2D) over vector clock (1D) and the significance of the additional information that is sent in matrix clock, do you think adding additional dimension (3D matrix and so on) to the clock will be useful? What more information can be sent by using additional dimensions?

Q2. Consider causal message ordering.

Assume vector clocks being used.

Consider an updation event at site D which was propagated to all sites at time  $T_{su}$ ; which A is aware of.

A wants to know if B got the update by looking at B's vector sent to A. Can A know?

Now A wants to know if C knows about the update by looking at B's vector. That is A is interested in  $V_C[D]$  but he needs to figure this out looking at B's vector in the msg from B to A.

Assume that C has sent B a msg recently where  $VB[C] \geq T_{su}$

Then is it true that C has received the update? Can A figure out if C has received the broadcast?

If the messages passed used matrix clocks instead of vector clocks then can A have answers to both problems above?